

(Chapter-5)

RSARIAL SEARCH

Yanmei Zheng

ADVERSARIAL SEARCH

- **Optimal decisions**
- **MinMax algorithm**
- **α - β pruning**
- **Imperfect, real-time decisions**



ADVERSARIAL SEARCH

➤ **Search problems seen so far:**

- ✓ Single agent.
- ✓ No interference from other agents and no competition.

➤ **Game playing:**

- ✓ Multi-agent environment.
- ✓ Cooperative games.
- ✓ Competitive games → **adversarial search**

➤ **Specifics of adversarial search:**

- ✓ Sequences of player's decisions we control.
- ✓ Decisions of other players we do not control.

Game

- Today's topic about games
- Two-player, turn-taking
- Fully observable, deterministic
- Zero-sum($1 + (-1) = 0$)
- Time-constrained

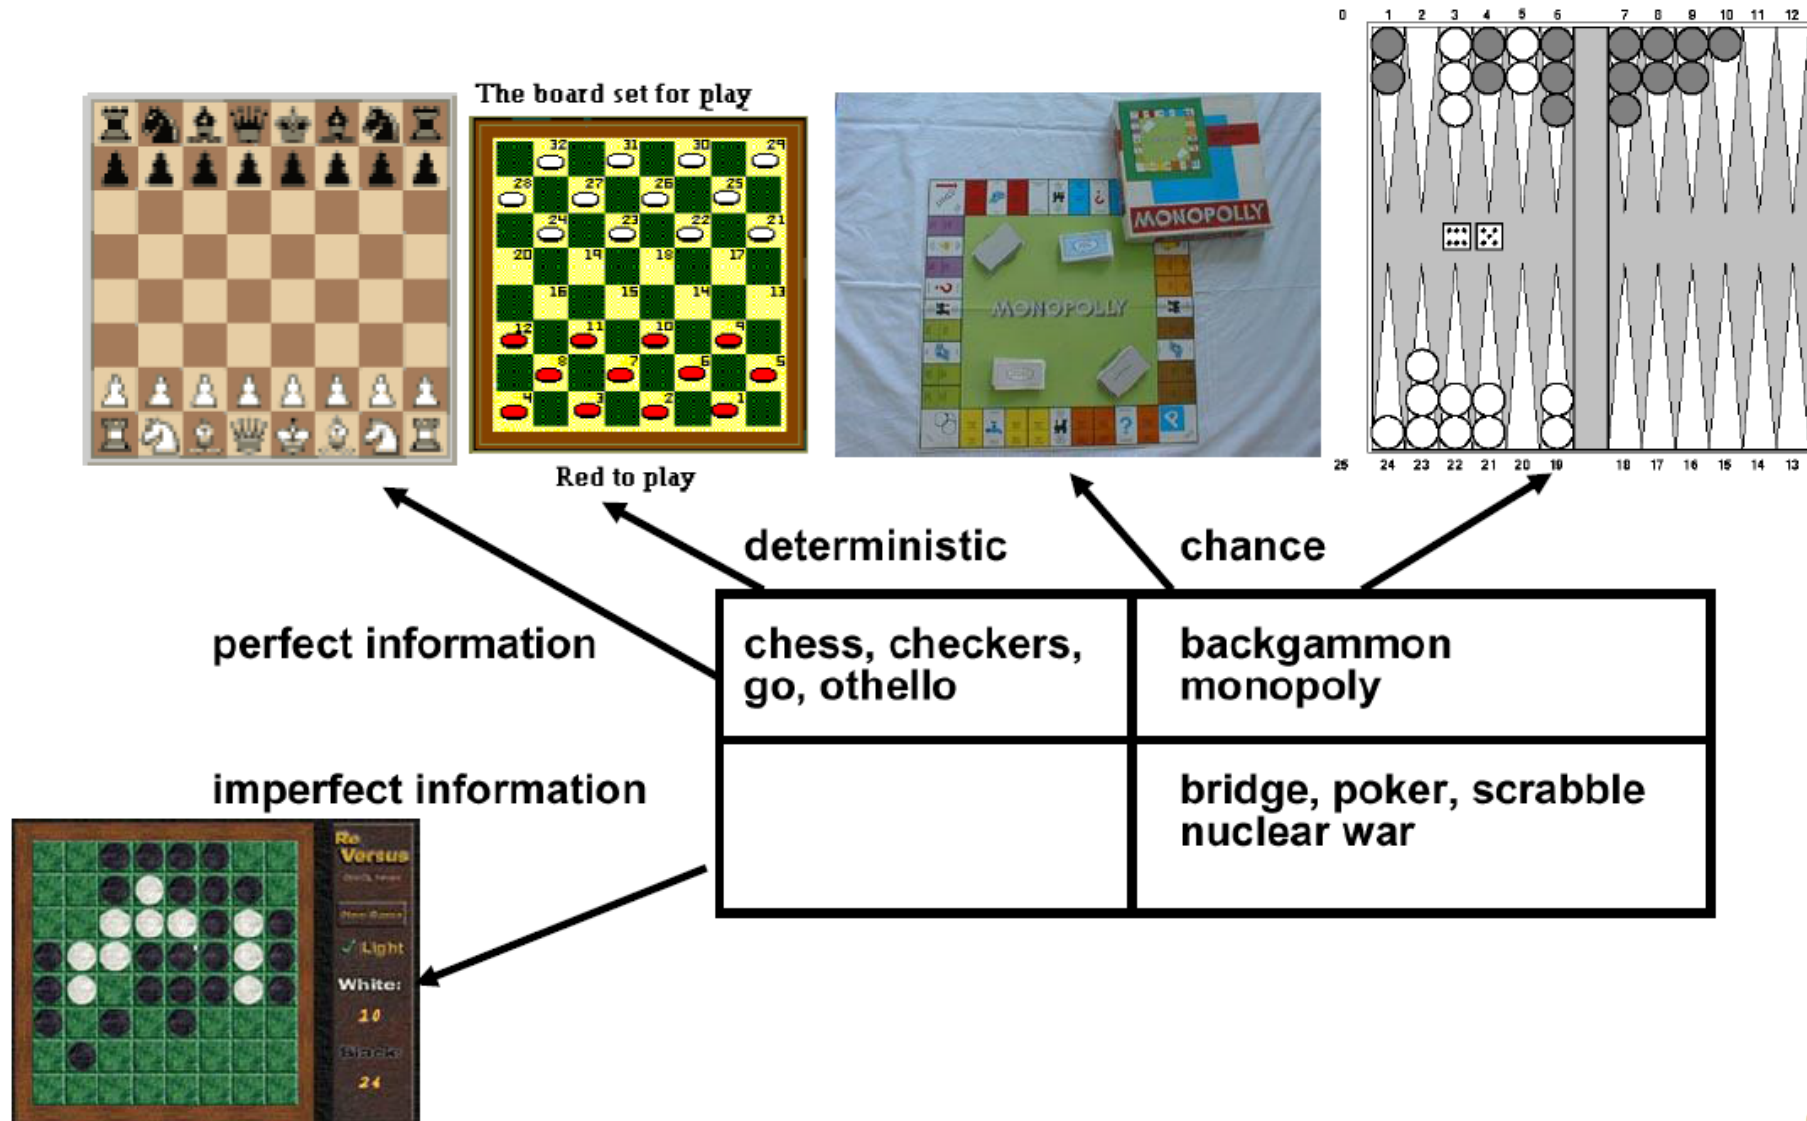
Games are a form of multi-agent environment

- **Multi-agent environment**
 - What do other agents do and how do they affect our success?
 - Cooperative vs. competitive multi-agent environments.
 - Competitive multi-agent environments give rise to adversarial search a.k.a. games

Why study games?


- **Why study games? ***
 - **Games are fun!**
 - **Historical role in AI**
 - **Studying games teaches us how to deal with other agents trying to foil our plans**
 - **Huge state spaces – Games are hard!**

Type of games



Game playing has a huge state space. How:

Chinese Chess



State space

Nine columns ten rows

Fourteen different kinds

Thirty-two pieces

Game playing has a huge state space.

- In general, the branching factor and the depth of terminal states are large.
 - Chess:
 - Number of states: $\sim 10^{40}$
 - Branching factor: ~ 35
 - Number of total moves in a game: ~ 100
 - Chinese chess、Go is more complicated
 - 10^{161} 、 10^{768}
- The chess search tree has about 35100 or 10154 nodes. If we want to use the complete search strategy, it would take an astronomical amount of time to process.
- The game requires a decision to be made even if the optimal decision cannot be calculated. How to make the best use of your time.
- The depth of the search tree affects performance.

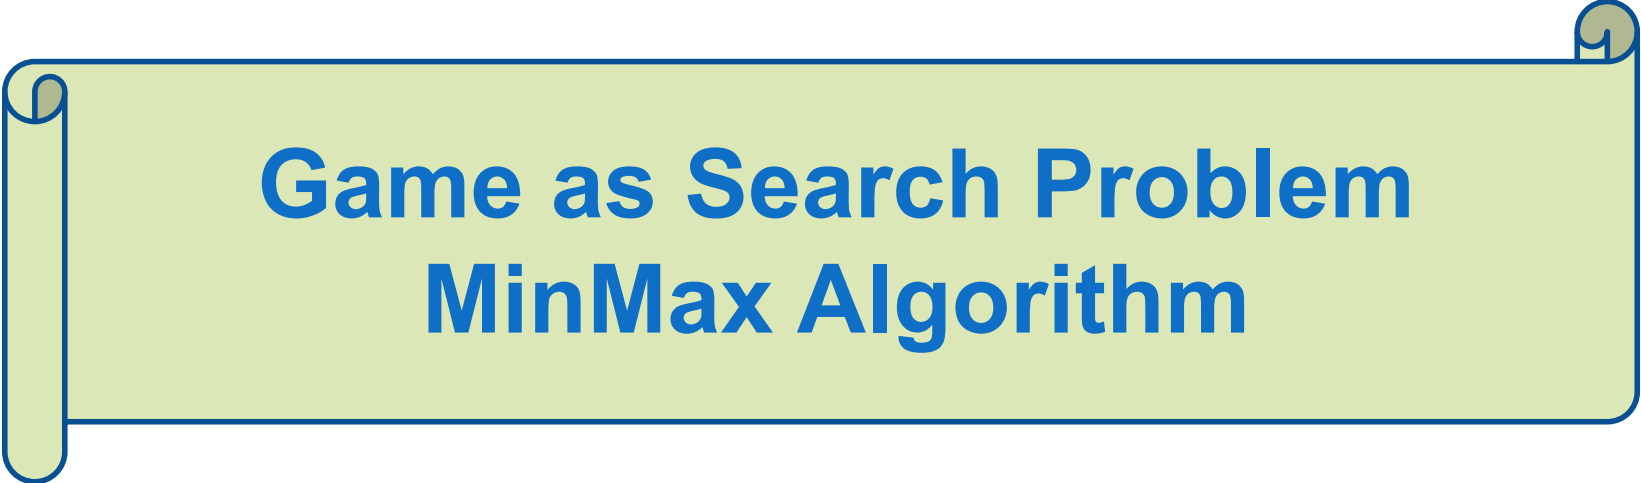
Games vs. search problems

- **"Unpredictable" opponent → specifying a move for every possible opponent reply**
- **Time limits → unlikely to find goal, must approximate**

-
- **How to deal with the huge state space?**

How to deal with the huge state space? (what are secrets?)



- Many game programs are based on
 - alpha-beta + iterative deepening + huge databases + ...
- The methods are general, but their implementation is dramatically improved by many specifically tuned-up enhancements (e.g., **the evaluation functions**).
- **Go** has too high a branching factor for existing search techniques. Current and future software must rely on **huge databases and pattern-recognition** techniques.
- **Search** is very important.



Game as Search Problem

MinMax Algorithm

Zero-Sum Games



- **Zero-Sum Games**

- Agents have opposite utilities (values on outcomes)
- Lets us think of a single value that one maximizes and the other minimizes
- Adversarial, pure competition

- **General Games**

- Agents have independent utilities (values on outcomes)
- Cooperation, indifference, competition, and more are all possible
- More later on non-zero-sum games

Two-Player Games


Two-Player Games with Complete Trees

- We can use search algorithms to write “intelligent” programs that **play games** against a human opponent.
- Just consider this extremely simple (and not very exciting) game:
 - At the beginning of the game, there are seven coins on a table.
 - Player 1 makes the first move, then player 2, then player 1 again, and so on.
 - One move consists of removing 1, 2, or 3 coins.
 - The player who removes all remaining coins wins.

Two-Player Games with Complete Trees

- Let us assume that the computer has the first move. Then, the game can be described as a **series of decisions**, where the first decision is made by the computer, the second one by the human, the third one by the computer, and so on, until all coins are gone.
- The **computer** wants to make decisions that **guarantee its victory** (in this simple game).
- The underlying assumption is that the **human** always finds the **optimal move**.


Two-Player Games with Complete Trees




Two-Player Games with Complete Trees

- So the computer will start the game by taking three coins and is **guaranteed to win** the game.
- The most practical way of implementing such an algorithm is the **Minimax procedure**:
 - Call the two players MIN and MAX.
 - Mark each leaf of the search tree with -1, if it shows a victory of MIN, and with 1, if it shows a victory of MAX.
 - Propagate these values up the tree using the rules:
 - If the parent state is a MAX node, give it the maximum value among its children.
 - If the parent state is a MIN node, give it the minimum value among its children.


Two-Player Games with Complete Trees



Adversarial Search



Tic-Tac-Toe Game Tree



Game setup

- Consider a game with Two players: **(Max)** and **(Min)**
- **Max** moves first and they take turns until the game is over.
Winner gets award, loser gets penalty.
- **Games as search:**
 - **Initial state:** e.g. board configuration of chess
 - **Successor function:** list of (move, state) pairs specifying legal moves.
 - **Goal test:** Is the game finished?
 - **Utility function:** Gives numerical value of terminal states. E.g. win (+1), lose (-1) and draw (0) in tic-tac-toe
- **Max** uses **search tree** to determine next move.

Example of an ADVERSARIAL two player Game

Tic-Tac-Toe (TTT)

- MAX has 9 possible first moves, etc.
- Utility value is always from the point of view of MAX.
- High values good for MAX and bad for MIN.

Two-Player Games with Complete Trees

- The previous example shows how we can use the Minimax procedure to determine the computer's best move.
- It also shows how we can apply depth-first search and a variant of backtracking to prune the search tree.
- Before we formalize the idea for pruning, let us move on to more interesting games of Tic-Tac-Toe.
- For such games, it is **impossible** to check every possible sequence of moves. The computer player then only looks ahead a certain number of moves and **estimates** the chance of winning after each possible sequence.

How to Play a Game by Searching

- **General Scheme**

- Consider all legal moves, each of which will lead to some new state of the environment ('board position')
- **Evaluate** each possible resulting board position
- Pick the move which leads to the best board position.
- Wait for your opponent's move, then **repeat**.

- **Key problems**

- Representing the 'board'
- Representing legal next boards
- Evaluating positions
- Looking ahead

Game Trees

- Represent the problem space for a game by a tree
 - ❖ **Nodes** represent 'board positions' (state)
 - ❖ **edges** represent legal moves.
- **Root node** is the position in which a decision must be made.
- **Evaluation function f** assigns real-number scores to 'board positions.'
- **Terminal nodes (leaf)** represent ways the game could end, labeled with the desirability of that ending (e.g. win/lose/draw or a numerical score)

MAX & MIN Nodes

- When **I** move, I attempt to **MAXimize** my performance.
- When my opponent moves, he attempts to **MINimize** my performance.


TO REPRESENT THIS:

- If we move first, label the root **MAX**; if our opponent does, label it **MIN**.
- **Alternate** labels for each successive tree level.
 - if the root (level 0) is our turn (MAX), all even levels will represent turns for us (MAX), and all odd ones turns for our opponent (MIN).


Evaluation functions

- Evaluations how good a 'board position' is
 - Based on **static features** of that board alone
- **Zero-sum** assumption lets us use **one function** to describe goodness for both players.
 - $f(n) > 0$ if we are winning in position n
 - $f(n) = 0$ if position n is tied
 - $f(n) < 0$ if our opponent is winning in position n
- Build using expert knowledge (**Heuristic**),
 - Tic-tac-toe: $f(n) = (\text{\# of 3 lengths possible for me}) - (\text{\# 3 lengths possible for you})$


Heuristic measuring for adversarial tic-tac-toe



X has 6 possible win paths:



O has 5 possible win:
 $E(n)=6-5=1$




Heuristic is $E(n)=M(n)-O(n)$

Where $M(n)$ is the total of My possible wining lines

$O(n)$ is total of Opponent's possible wining lines

$E(n)$ is the total Evaluation for state n

Heuristic measuring for adversarial tic-tac-toe



X has 4 possible win paths;

O has 6 possible wins

$$E(n)=4-6=-2$$


Heuristic is $E(n)=M(n)-O(n)$

Where $M(n)$ is the total of My possible wining lines

$O(n)$ is total of Opponent's possible wining lines

$E(n)$ is the total Evaluation for state n

Heuristic measuring for adversarial tic-tac-toe



X has 5 possible win paths;

O has 4 possible wins

$$E(n) = 5 - 4 = 1$$

Heuristic is $E(n) = M(n) - O(n)$

Where $M(n)$ is the total of My possible winning lines

$O(n)$ is total of Opponent's possible winning lines

$E(n)$ is the total Evaluation for state n

Maximize $E(n)$ $E(n) = 0$ when my opponent and I have equal number of possibilities.

Two-Player Games

$$E(n) = 8 - 8 = 0$$

X		
O	X	

$$E(n) = 6 - 2 = 4$$

O	O	X
X	O	
X		

$$E(n) = 2 - 2 = 0$$

shows the weak-ness of
this $e(p)$

How about these?

O	O	X
	X	
X		

$$E(n) =$$

X	X	
O	O	O
	X	


$$E(n) = -$$

MinMax Algorithm

- **Main idea:** choose move to position with highest minimax value. = best achievable payoff against best play.
- E.g., 2-ply game:

Adversarial Search (Minimax)

- Deterministic, zero-sum games:
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- Minimax search:
 - A state-space search tree
 - Players alternate turns
 - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



Optimal decision in games{5.2}





Figure 5.2 A two-ply game tree. The Δ nodes are “MAX nodes,” in which it is MAX’s turn to move, and the ∇ nodes are “MIN nodes”. The terminal nodes show the utility values for MAX; the other nodes are labeled with their minimax values. MAX’s best move at the root is a_1 , because it leads to the state with the highest minimax value, and MIN’s best reply is b_1 , because it leads to the state with the lowest minimax value.

Optimal decision in games{5.2}



Minmax(s)=

$$\left\{ \begin{array}{ll} \text{UTILITY}(S) & \text{if } \text{TERMINAL} - \text{TEST}(s) \\ \max_{a \in \text{Action}(s)} \text{MINMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Action}(s)} \text{MINMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{array} \right\}$$

MinMax Algorithm


function MINIMAX-DECISION(*state*) returns an *action*
 $v \leftarrow \text{MAX-VALUE}(\textit{state})$
 return the action in **SUCCESSORS(*state*)** with value *v*

function MAX-VALUE(*state*) returns a *utility value*
 if **TERMINAL-TEST(*state*)** then return **UTILITY(*state*)**
 $v \leftarrow -\infty$
 for *a, s* in **SUCCESSORS(*state*)** do
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$
 return *v*

function MIN-VALUE(*state*) returns a *utility value*
 if **TERMINAL-TEST(*state*)** then return **UTILITY(*state*)**
 $v \leftarrow -\infty$
 for *a, s* in **SUCCESSORS(*state*)** do
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$
 return *v*

Minimax Implementation

```
def max-value(state):  
    initialize v =  $-\infty$   
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```



$$V(s) = \max_{s' \in \text{successor}(s)} V(s')$$

```
def min-value(state):  
    initialize v =  $+\infty$   
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s') = \min_{s \in \text{successor}(s')} V(s)$$

Minimax Implementation (Dispatch)

def value(state):

if the state is a terminal state: return the state's utility

if the next agent is **MAX**: return **max-value(state)**

if the next agent is **MIN**: return **min-value(state)**

def max-value(state):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

def min-value(state):

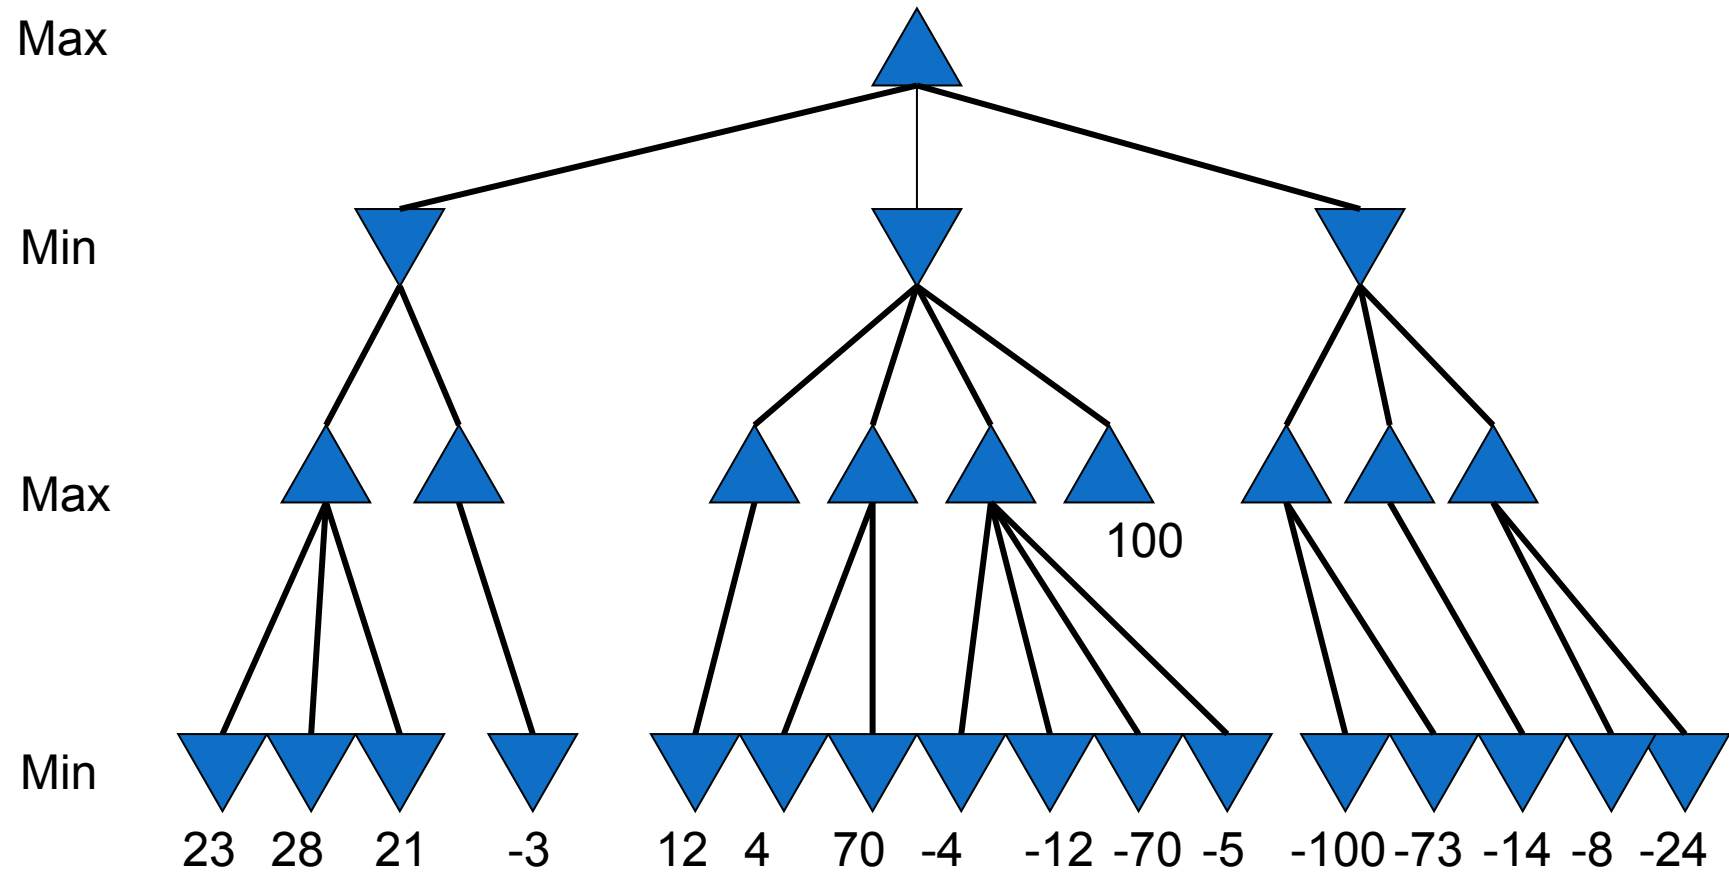
initialize $v = +\infty$

for each successor of state:

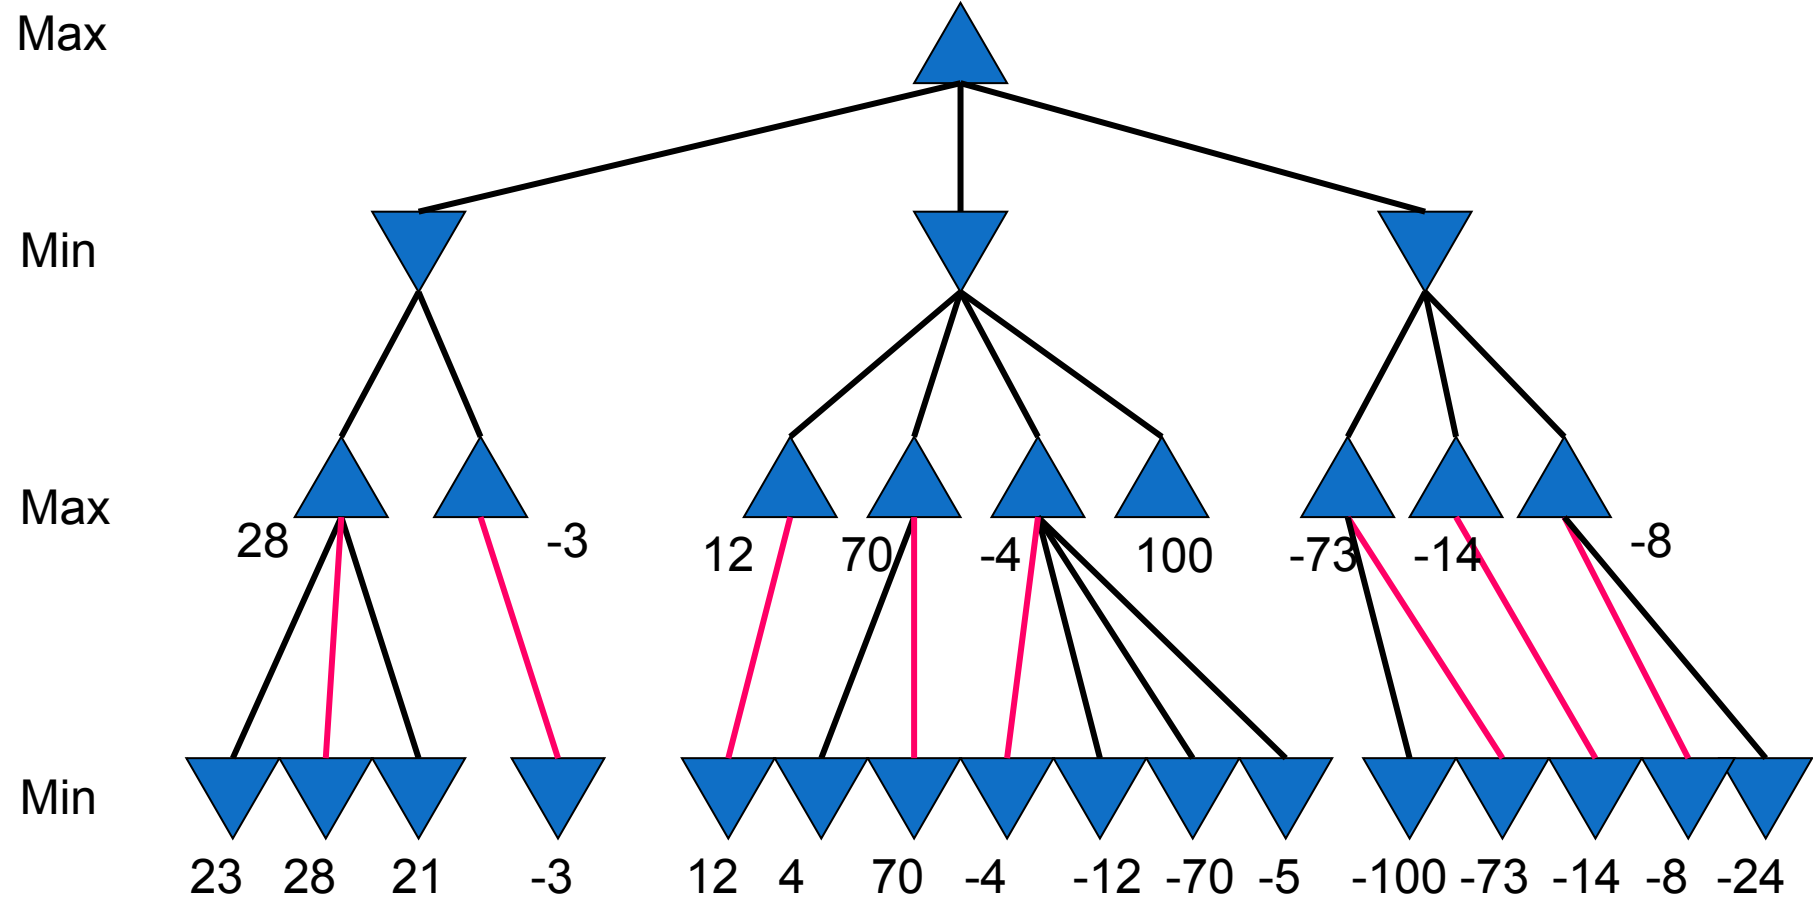
$v = \min(v, \text{value}(\text{successor}))$

return v


Minimax tree




Minimax tree




Minimax tree



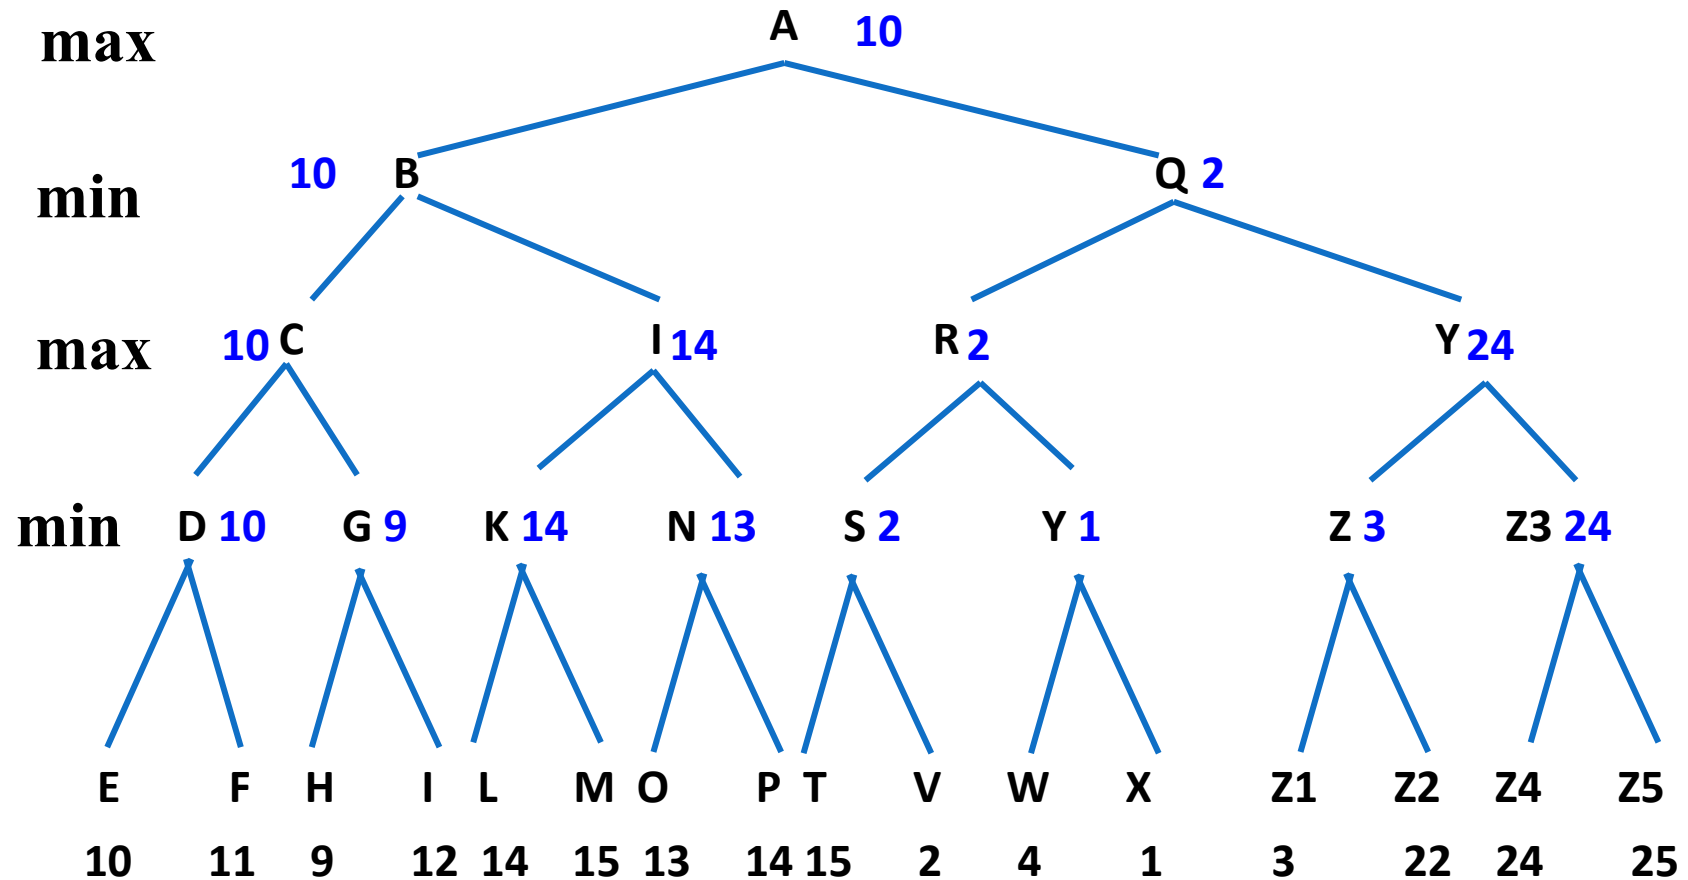
Minimax tree




Minimax tree



Minimax tree




Game Tree Pruning



Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?




MinMax Analysis

- Time Complexity: $O(b^d)$
- Space Complexity: $O(b*d)$
- Optimality: Yes




Problem: Game → Resources Limited!

- Time to make an action is limited
- Can we do better ? Yes !
- How ? Cutting useless branches !


Some nodes in the search can be *proven* to be irrelevant to the outcome of the search



Minimax Properties




Optimal against a perfect player. Otherwise?



$$\text{MINMAX}(\text{root}) = \max(\min(3, 12, 8), \min(2, 4, 6), \min(14, 5, 2))$$


Which nodes needn't to be searched?
Which value dose the result of Max () is independent of?

Alpha-Beta Pruning{5,3}




$$\begin{aligned}\text{MINMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\ &= 3\end{aligned}$$

Minimax Example




Minimax Pruning



Alpha-Beta Pruning

- **General configuration (MIN version)**
 - We're computing the MIN-VALUE at some node n
 - We're looping over n 's children
 - n 's estimate of the childrens' min is dropping
 - Who cares about n 's value? MAX
 - Let a be the best value that MAX can get at any choice point along the current path from the root
 - If n becomes worse than a , MAX will avoid it, so we can stop considering n 's other children (it's already bad enough that it won't be played)
- **MAX version is symmetric**



The Alpha-Beta Procedure


- Now let us specify how to prune the Minimax tree in the case of a static evaluation function.
- Use two variables alpha (associated with MAX nodes) and beta (associated with MIN nodes).
- These variables contain the best (highest or lowest, resp.) $E(p)$ value at a node p that has been found so far.
- Notice that alpha can never decrease, and beta can never increase.

The Alpha-Beta Procedure


- There are two rules for terminating search:
 - Search can be stopped below any MIN node having a beta value less than or equal to the alpha value of any of its MAX ancestors.
 - Search can be stopped below any MAX node having an alpha value greater than or equal to the beta value of any of its MIN ancestors.
- Alpha-beta pruning thus expresses a relation between nodes at level n and level $n+2$ under which entire subtrees rooted at level $n+1$ can be eliminated from consideration.

α Cuts


- If the **current max** value is greater than the successor's min value, **don't** explore that min subtree any more



α Cut example




α Cut example




- Depth first search along path 1

α Cut example




- 21 is minimum so far (second level)
- Can't evaluate yet at top level

α Cut example




- -3 is minimum so far (second level)
- -3 is maximum so far (top level)

α Cut example




- -70 is now minimum so far (second level)
- -3 is still maximum (can't use second node yet)

α Cut example




- Since second level node will never be > -70 , it will never be chosen by the previous level
- We can stop exploring that node

α Cut example




- Evaluation at second level is -73

α Cut example



- Again, can apply α cut since the second level node will never be > -73 , and thus will never be chosen by the previous level

α Cut example




- As a result, we evaluated the Max node without evaluating several of the possible paths

β cuts

- Similar idea to α cuts, but the other way around
- If the current minimum is less than the successor's max value, don't look down that max tree any more

β Cut example



- Some subtrees at second level already have values $>$ min from previous, so we can stop evaluating them.

α - β Pruning

- **Pruning by these cuts does not affect final result**
 - May allow you to go much deeper in tree
- “Good” ordering of moves can make this pruning much more efficient
 - Evaluating “best” branch first yields better likelihood of pruning later branches
 - Perfect ordering reduces time to $b^{m/2}$
 - i.e. doubles the depth you can search to!

Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root

def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$ return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \leq \alpha$ return v

$\beta = \min(\beta, v)$

 return v

α - β Pruning

- Can store information along an entire *path*, not just at most recent levels!
- Keep along the path:
 - α : **best MAX** value found on this path
(initialize to most negative utility value)
 - β : **best MIN** value found on this path
(initialize to most positive utility value)





The Alpha and the Beta

- For a leaf, $\alpha = \beta = \text{utility}$
- At a max node:
 - $\alpha = \text{largest child utility found so far}$
 - $\beta = \beta \text{ of parent}$
- At a min node:
 - $\alpha = \alpha \text{ of parent}$
 - $\beta = \text{smallest child utility found so far}$
- For any node:
 - $\alpha \leq \text{utility} \leq \beta$
 - “If I had to decide now, it would be...”







Alpha-Beta example

Alpha-Beta Pruning {5,3}



Each node is marked with a **range** of its value

Alpha-Beta Pruning {5,3}




Each node is marked with a **range** of its value

Alpha-Beta Pruning {5.3}

MAX


$\alpha = -\infty, \beta = +\infty$

MIN



- ❑ α = The best value (or maximum) of MAX found on the path so far
- ❑ β = The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}



- ❑ α = The best value (or maximum) of MAX found on the path so far
- ❑ β = The best value (or minimum) of MIN found on the path so far


Alpha-Beta Pruning {5.3}

MAX

$\alpha = -\infty, \beta = +\infty$

MIN

$\alpha = -\infty, \beta = 3$




- ❑ α = The best value (or maximum) of MAX found on the path so far
- ❑ β = The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}

MAX

$\alpha = -\infty, \beta = +\infty$

MIN




- ❑ α = The best value (or maximum) of MAX found on the path so far
- ❑ β = The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}

MAX


$\alpha = -\infty, \beta = +\infty$

MIN




- ❑ α = The best value (or maximum) of MAX found on the path so far
- ❑ β = The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- ❑ α =The best value (or maximum) of MAX found on the path so far
- ❑ β =The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- ❑ α = The best value (or maximum) of MAX found on the path so far
- ❑ β = The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- ❑ α =The best value (or maximum) of MAX found on the path so far
- ❑ β =The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- α =The best value (or maximum) of MAX found on the path so far
- β =The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- α =The best value (or maximum) of MAX found on the path so far
- β =The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- α =The best value (or maximum) of MAX found on the path so far
- β =The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- α =The best value (or maximum) of MAX found on the path so far
- β =The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}




- α =The best value (or maximum) of MAX found on the path so far
- β =The best value (or minimum) of MIN found on the path so far

Alpha-Beta Pruning {5.3}



- α =The best value (or maximum) of MAX found on the path so far
- β =The best value (or minimum) of MIN found on the path so far


Alpha-Beta Pruning {5.3}



$$\alpha = -\infty, \beta = 3$$


The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:



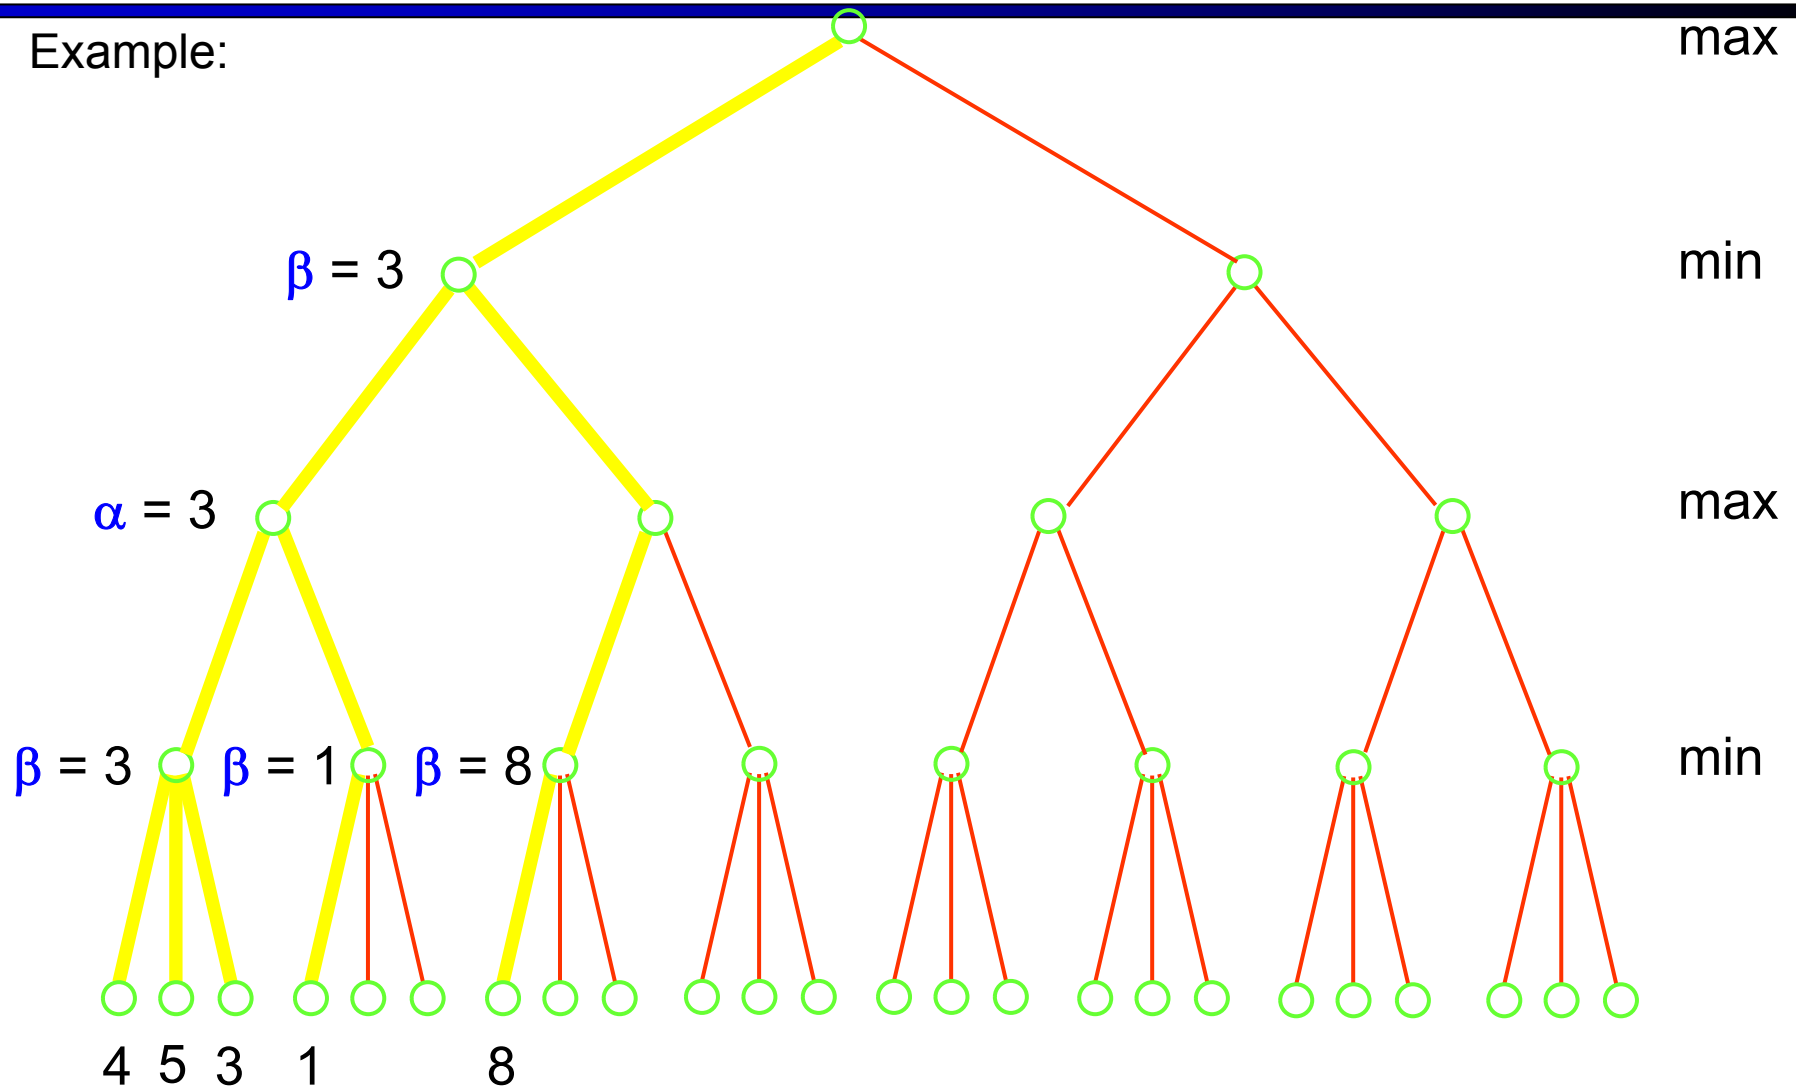
The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure



The Alpha-Beta Procedure

Example:


$\alpha = 3$



The Alpha-Beta Procedure

Example:


$\alpha = 3$



The Alpha-Beta Procedure


Example:

$\alpha = 3$




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure

Example:




The Alpha-Beta Procedure


Example:




Alpha-Beta Quiz 1



Alpha-Beta Quiz2




Alpha-Beta Quiz 3



Move ordering{5.3}


MAX



if changing the ordering: 2 5 14

Move ordering{5.3}

MAX




if changing the ordering: 2 5 14

It is better if the **MAX children of a MIN node** are ordered in **increasing** backed up values

Move ordering {5.3.1}


- ❑ Find the best move:
- ❑ Which nodes are pruned?



The MAX children of MIN nodes are ordered in increasing back up values

Move ordering {5.3.1}

- ❑ Find the best move:
- ❑ Which nodes are pruned?



It is better if the **MIN children of a MAX node** are ordered in **decreasing** backed up values

Move ordering {5.3.1}

- Assume a game tree of uniform branching factor b
- Minimax examines $O(b^m)$ nodes, so does alpha-beta in the worst-case
- The gain for alpha-beta is **maximum** when:
 - ✓ The **MIN children of a MAX** node are ordered in **decreasing** backed up values
 - ✓ The **MAX children of a MIN** node are ordered in **increasing** backed up values
 - ✓ then alpha-beta examines $O(b^{m/2})$ nodes
- But this requires an oracle (if we knew how to order nodes perfectly, we would

_____)

Move ordering {5.3.1}

- Assume a game tree of uniform branching factor b
- Minimax examines $O(b^m)$ nodes, so does alpha-beta in the worst-case
- The gain for alpha-beta is **maximum** when:
 - ✓ The **MIN children of a MAX** node are ordered in **decreasing** backed up values
 - ✓ The **MAX children of a MIN** node are ordered in **increasing** backed up values
 - ✓ then alpha-beta examines $O(b^{m/2})$ nodes
- But this requires an oracle (if we knew how to order nodes perfectly, we would **not need to search the tree**)
- If nodes are ordered at random, then the average number of nodes examined by alpha-beta is $\sim O(b^{3m/4})$

-
- **We often can not reach to the terminal nodes within limited time!**
 1. **When search to the limited depth, just cutoff and replace the Utility() with EVAL()**
 2. **Iterative deepening.**

Imperfect real-time decisions {5.4}

- Search should be truncated early
- Replace the utility function with the heuristic evaluation function EVAL, which estimates the utility value of a board game, and the decision when to use an EVAL cutoff test instead of terminating the test

H-MINMAX(s, d)=

$$\left\{ \begin{array}{ll} \text{EVAL}(S) & \text{if } \text{CUTOFF} - \text{TEST}(s, d) \\ \max_{a \in \text{Action}(s)} H - \text{MINMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Action}(s)} H - \text{MINMAX}(\text{RESULT}(s, a), d + 1) & \text{if } \text{PLAYER}(s) = \text{MIN}. \end{array} \right\}$$


MINMAX(s)=

$$\left\{ \begin{array}{ll} \text{UTILITY}(S) & \text{if } \text{TERMINAL} - \text{TEST}(s) \\ \max_{a \in \text{Action}(s)} \text{MINMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Action}(s)} \text{MINMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{array} \right\}$$

Evaluation functions {5.4.1}

- It is obvious that the performance of a game program depends heavily on the quality of the evaluation function.

Practice




node	B	C	D	E	F	G
(\uparrow \uparrow) obtained from the parent						
Value returned						

Practice:

MAX

MIN

MAX



node	B	C	D	E	F	G
(\lceil \lceil) obtained from the parent	$-\infty, +\infty$	3, $+\infty$	$-\infty, +\infty$	$-\infty, 3$	3, $+\infty$	Null
Value returned	3	3	3	4	3	Null



Thank you

**End of
Chapter 5**